

CalmGuide — Weekly Report

Reporting period: 2026-09-18 → 2026-09-24 **Prepared for:** Leap of Faith (LOF) weekly deliverable

Repository state: main @ b277688

1. Summary

This week's work fell into three areas: a full visual redesign of the web and mobile applications, a set of functional defect fixes in the caregiver dashboard, and the discovery and repair of a deployment failure that had been silently blocking releases for twelve days.

Nine commits landed, touching 125 files (+3,515 / -801 lines) across the web frontend, mobile app, backend and project documentation. A new Android build was produced and is available for testing.

The most consequential item is not a feature. The live web application had not received a successful deployment since 2026-09-08. This was found, diagnosed and fixed this week, and is documented in §4.

2. What Shipped

2.1 Visual design system

A new design direction was applied across the web application and carried into the mobile theme: a near-black ink, sky-blue and coral palette over pale atmospheric gradients, with translucent "glass" panels on primary content surfaces.

Two deliberate constraints were built in:

uniformly. A heavy backdrop blur on a small control costs the same rendering work while reducing legibility.

reports a reduced-transparency preference, and when the browser does not support backdrop filtering. Reduced transparency is a genuine accessibility setting among older users, who are a substantial part of the caregiver population.

- The glass treatment is applied to two large content panels only, not
- The material degrades to solid surfaces when the operating system

2.2 Dashboard restructure

profile and the two primary actions.

element on the screen, so attention lands there first.

- Two-column desktop layout with a persistent left rail carrying the care
- Moment Coach presented as the single dark surface and the only animated
- A weekly progress chart driven by real check-in data.

2.3 Profile portraits

Caregivers may now choose an illustrated portrait for the person they care for, selected during signup alongside the name.

This is an explicit choice, not an inference. The backend stores no name and no gender by design, so nothing server-side could determine a portrait, and deriving one from a name would misgender people. A lettered monogram remains a first-class third option rather than a fallback.

3. Defects Fixed

3.1 Weekly progress chart could never display a trend

The chart was built against an endpoint that filters to a single calendar date, so it was structurally incapable of plotting more than one point. A read-only history endpoint was added to the backend.

A second defect in the same feature: dates returned as YYYY-MM-DD were parsed as UTC midnight and then read in local time, shifting every entry one day earlier in any timezone west of UTC and dropping the current day off the chart entirely. Both are fixed and covered by tests that run under four timezones.

3.2 Sign-out did not sign out

Signing out cleared the session keys but left the stored profile list behind. The application then treated that surviving profile as a session to resume, rebuilt the credentials from it, and returned the user to the dashboard. Sign-out now clears the full session.

3.3 Confirmation dialog rendered behind page content

The sign-out confirmation appeared with its heading obscured and dashboard panels painted over it. The dialog relied on escaping to the root stacking context; a layout change introduced earlier in the week created an intervening stacking context that trapped it. The dialog is now rendered through a portal, which makes it independent of the surrounding layout.

4. Deployment Incident

Impact: the production web application served a twelve-day-old build while the codebase continued to advance. Any review of the live site during that period was not looking at current work.

Cause: the hosting provider refuses a deployment when the GitHub account that authored the commit does not hold a seat on the project team. The build is skipped entirely — no compilation is attempted and no build log is produced. This began when the hosting plan changed on 2026-09-07.

Why it went unnoticed: the failure is silent by design. The dashboard reports a generic failure with no log, the previous build continues to serve normally, and nothing distinguishes a stale site from a healthy one.

Resolution: commit authorship was corrected to an account holding a seat. Verified on both production and preview deployments, which now build and publish normally. The failure mode and its fix are documented in the repository README so it cannot recur unrecognised.

5. Guidance Quality Review

A CalmGuide response was compared against a general-purpose assistant's answer to the same caregiving scenario — a person with middle-stage dementia repeatedly asking the same question.

Three defects were traced to specific instructions in our own prompt:

guidance lists as a mistake to avoid during a crisis. The prohibition had been written as advice about the caregiver's instincts, so it never constrained the assistant's own suggested wording.

redirected. For repeated questioning, the answer itself carries most of the reassurance.

Repetitive questioning is a common proxy for toilet, thirst, hunger, pain, temperature or fatigue, and a person in middle-stage dementia may no longer be able to name discomfort directly.

- The response scripted the phrase "let me explain", which our own
- The response never answered the caregiver's question — it deflected and
- The response never suggested checking for an unmet physical need.

All three were addressed and confirmed against live model generations rather than by inspection of the prompt alone. One behaviour — the unmet needs check — currently appears in approximately half of generations and is being made deterministic.

An earlier attempt to mandate emergency red-flag signs in every response was tried and withdrawn: the model did not follow the instruction at two prompt positions, and the application already presents emergency contacts persistently on screen, so inlining them on routine questions adds alarm without adding reach.

6. Testing

Suite	Result
Backend (pytest)	1,837 passed
Web frontend (vitest)	276 passed, 1 skipped
Mobile (jest)	28 passed
Type checking	Clean across web and mobile
Production build	Passes

A prompt cost-budget test correctly failed during this week's prompt work and forced a deliberate decision on token growth rather than allowing it to drift. The budget was raised from 5,000 to 5,150 tokens per request (measured 4,913 → 5,082, +3.4%) with the reasoning recorded in the test.

7. Known Limitations and Open Items

against the production environment**.

applies inconsistently.

chart, sign-out fix and dialog fix are web-only; the mobile application has a separate storage layer and screen set, and none of this work was ported.

credentials, which caused avoidable confusion this week.

- The coach prompt changes are verified locally but **not yet verified
- The unmet-needs check in guidance responses is non-deterministic and
- **Mobile received the palette and icons only.** The portraits, progress
- Access-code documentation does not distinguish local from production

8. Planned — Remainder of Sprint

OpenEMR integration for the patient profile (feasibility spike).

An optional enrichment layer: a caregiver creates a CalmGuide profile as today, and may then connect an existing OpenEMR record. CalmGuide retains orchestration — the external record is a data source, not something that addresses the model directly. Patient clinical context is kept separate from the dementia-care knowledge base, safety routing is unchanged, and the service layer is provider-agnostic with OpenEMR as the first implementation. Initial scope: Patient, Condition, AllergyIntolerance, Medication, Observation, Encounter. Authorization via OAuth with tokens encrypted at rest.

This is time-boxed with a go/no-go recommendation at sprint end. Two items require LOF and clinical input before it becomes a committed feature: the available FHIR surface varies by OpenEMR deployment and version, and the profile model currently holds clinical data and no personally identifying information, so connecting a medical record changes the project's data-handling posture.

Psychology assessment form. Shared with Chaitanya Chennupati (M.Sc. Neuropsychology, Christ University, Bangalore) for review, with backend integration to follow so assessment data can inform the care profile and guidance.

9. Deliverables

Item	Status
Weekly report	This document
Repository archive	Provided (source and documentation, credentials excluded)
Video demonstration	To follow